

<https://github.com/ISE-FIZKarlsruhe/bikidata>

# bikidata

FIZ-ISE Seminar 16 Aug 2023

Etienne Posthumus

[eposthumus@gmail.com](mailto:eposthumus@gmail.com)    <https://epoz.org/>

bietjie = beetje

“A little bit”

*pronounced:* biki

# Why?

We often use Wikidata in our research.

...for queries

...for extraction

But the WDQS has time-outs, or we want to do something more advanced for analysis

*Can we share efforts?* Instead of all doing custom grep etc. separately

# Own triplestore?

Blazegraph / Virtuoso / QLever etc.

This is a non-trivial and time-consuming exercise.

(never mind resource-intensive)

# TL;DR

- Download Truthy Wikidata .nt dump, split into separate chunks.
- Hash, s-p-o IRIs to 64-bit integers
- Hash IRI/Literal to 64-bit integers
- Store everything in Parquet files
- Query using DuckDB
- ...
- 🎉

# Choosing a hash algorithm

Murmurhash, SIP, Cityhash, Spookyhash, xxhash, FNV, MD5, CRC, SHA

And does it really need to be 64-bits? Isn't 32 bits enough?

# Choosing a hash algorithm

Murmurhash, SIP, Cityhash, Spookyhash, xxhash, FNV, MD5, CRC, SHA

And does it really need to be 64-bits? Isn't 32 bits enough?

...and the winner is: **xxhash**

<https://github.com/Cyan4973/xxHash>

# Looks like

s	p	o
10504872225212057570	12580884072737042117	1877790563993361007
10504872225212057570	1969734015126054522	5587686886852776007
10504872225212057570	15314476376240055071	5464438232649063526
10504872225212057570	3176906616339849618	8616764692071278374
10504872225212057570	3176906616339849618	8661715247485442335
10504872225212057570	3176906616339849618	1734327773552895906
10504872225212057570	3176906616339849618	11573360774852176845



# Looks like

s	p	o
10504872225212057570	12580884072737042117	1877790563993361007
10504872225212057570	1969734015126054522	5587686886852776007
10504872225212057570	15314476376240055071	5464438232649063526
10504872225212057570	3176906616339849618	8616764692071278374
10504872225212057570	3176906616339849618	8661715247485442335
10504872225212057570	3176906616339849618	1734327773552895906
10504872225212057570	3176906616339849618	11573360774852176845

10504872225212057570 | <<http://www.wikidata.org/entity/Q42>>

# Looks like

10504872225212057570	13787943592591698645	18303551757152610549	"4ce2ba117755a"
10504872225212057570	13149575899394925943	5620510158999179407	"39a33dc4-5a81-4d67-91d6-1daecdb854e3"
10504872225212057570	16872716867785849655	9380107508382309212	"2078791"
10504872225212057570	11643480590671441163	16342048489805386123	<http://www.wikidata.org/entity/Q1860>
10504872225212057570	17591124141766173371	11286963682984007955	"ncf10168152"
10504872225212057570	2502498544802895930	14455623418673618869	"DouglasAdams"
10504872225212057570	12928713851174965354	11645328310747403540	<http://kbpedia.org/kko/rc/DouglasAdams>
10504872225212057570	12457161947493164903	17795690644028887794	"26792807"
10504872225212057570	4881929489000864687	2501072958756939372	<http://www.wikidata.org/entity/Q2687578>
10504872225212057570	17657798667172610268	4885215682934094547	<http://id.loc.gov/authorities/classification/PR6051.D3352>
10504872225212057570	3029897129149013315	9225015895394000726	"215957"
10504872225212057570	4497732400159369300	18099947938301090450	<http://www.wikidata.org/entity/Q6173448>
10504872225212057570	2002846622502823662	13975239749449742978	<http://www.wikidata.org/entity/Q7066>
10504872225212057570	9838106236933647348	11690699215486955192	"DNA"@en

# DuckDB

<https://duckdb.org/>

An in-process SQL OLAP database management system

# Parquet files

<https://parquet.apache.org/>

Apache Parquet is an open source, column-oriented data file format designed for efficient data storage and retrieval. It provides efficient data compression and encoding schemes with enhanced performance to handle complex data in bulk.

# Put it all together and we can...

```
select count(*) from 'tri.parquet'
```

```
where p=1969734015126054522 and o = 5587686886852776007
```

```
10 646 247
```

# Put it all together and we can...

```
select count(*) from 'tri.parquet'
```

```
where p=1969734015126054522 and o = 5587686886852776007
```

```
10 646 247
```

```
select hash, literal from 'map.parquet' where hash in  
(1969734015126054522, 5587686886852776007)
```

```
1969734015126054522 | <http://www.wikidata.org/prop/direct/P31>
```

```
5587686886852776007 | <http://www.wikidata.org/entity/Q5>
```

Or more usefully:

```
WITH Q53592_po AS (SELECT p,o FROM 'xa?.parquet' WHERE s = 12746726515823639617)
SELECT p_cnt, (SELECT iri FROM 'index.parquet' WHERE hash = s)
  FROM (SELECT t.s, count(t.p) p_cnt FROM 'xa?.parquet' t
        INNER JOIN Q53592_po ON t.p = Q53592_po.p AND t.o = Q53592_po.o
        GROUP BY t.s
        ORDER BY count(t.p) DESC)
WHERE p_cnt > 10;
```





<http://www.wikidata.org/entity/Q13442814>	39371483
<http://www.wikidata.org/entity/Q5>	10646247
<http://www.wikidata.org/entity/Q4167836>	5135154
<http://www.wikidata.org/entity/Q16521>	3552731
<http://www.wikidata.org/entity/Q523>	3291341
<http://www.wikidata.org/entity/Q7318358>	2096006
<http://www.wikidata.org/entity/Q318>	2091076
<http://www.wikidata.org/entity/Q4167410>	1436977
<http://www.wikidata.org/entity/Q113145171>	1252930
<http://www.wikidata.org/entity/Q11173>	1249605
<http://www.wikidata.org/entity/Q7187>	1213201
<http://www.wikidata.org/entity/Q8054>	992097
<http://www.wikidata.org/entity/Q11266439>	797866
<http://www.wikidata.org/entity/Q3305213>	695859
<http://www.wikidata.org/entity/Q79007>	650505

```
create temp table tipes as select o,
count(s) t from 'tri.parquet' where
p=1969734015126054522 group by o;

select m.literal, tipes.t from tipes join
'map.parquet' m on tipes.o = m.hash
order by t desc limit 15;
```

# Try it

teach02:/home/wikidata/

42M duckdb (the executable)

17G map.parquet

41G tri.parquet

# Try it

```
teach02:/home/wikidata/
```

```
pip install duckdb
```

```
https://github.com/ISE-FIZKarlsruhe/bikidata
```

# TODO list

- Want to sort the triple table, but how? Does that improve query speed?

# TODO list

- Want to sort the triple table, but how? Does that improve query speed?
- Make smaller extracts:
  - Persons (P31 Q5)
  - Person-plus1
  - Books
  - Which others?*

# TODO list

- Want to sort the triple table, but how? Does that improve query speed?
- Make smaller extracts:
  - Persons (P31 Q5)
  - Person-plus1
  - Books
  - Which others?*
- A cookbook of sample queries

# TODO list

- Want to sort the triple table, but how? Does that improve query speed?
- Make smaller extracts:
  - Persons (P31 Q5)
  - Person-plus1
  - Books
  - Which others?*
- A cookbook of sample queries
- A “labels” service, similar to the WDQS SERVICE

# TODO list

- Want to sort the triple table, but how? Does that improve query speed?
- Make smaller extracts:
  - Persons (P31 Q5)
  - Person-plus1
  - Books
  - Which others?*
- A cookbook of sample queries
- A “labels” service, similar to the WDQS SERVICE
- Better Python library, publish to PYPI



# TODO list

- Want to sort the triple table, but how? Does that improve query speed?
- Make smaller extracts:
  - Persons (P31 Q5)
  - Person-plus1
  - Books
  - Which others?*
- A cookbook of sample queries
- A “labels” service, similar to the WDQS SERVICE
- Better Python library, publish to PYPI

# TODO list

- Want to sort the triple table, but how? Does that improve query speed?
- Make smaller extracts:
  - Persons (P31 Q5)
  - Person-plus1
  - Books
  - Which others?*
- A cookbook of sample queries
- A “labels” service, similar to the WDQS SERVICE
- Better Python library, publish to PYPI
- Index literals using embeddings and a HNSW (SBERT + FAIS?)

# TODO list

- Want to sort the triple table, but how? Does that improve query speed?
- Make smaller extracts:
  - Persons (P31 Q5)
  - Person-plus1
  - Books
  - Which others?*
- A cookbook of sample queries
- A “labels” service, similar to the WDQS SERVICE
- Better Python library, publish to PYPI
- Index literals using embeddings and a HNSW (SBERT + FAIS?)

*Your suggestions?*